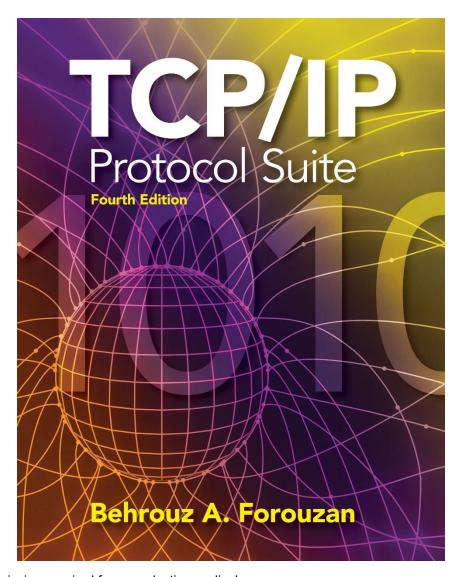
The McGraw·Hill Companies

Chapter 9

Internet
Control
Message
Protocol
Version 4 (ICMPV4)



OBJECTIVES:

- **□** To discuss the rationale for the existence of ICMP.
- To show how ICMP messages are divided into two categories: error reporting and query messages.
- To discuss the purpose and format of error-reporting messages.
- **□** To discuss the purpose and format of query messages.
- To show how the checksum is calculated for an ICMP message.
- To show how debugging tools using the ICMP protocol.
- ☐ To show how a simple software package that implements ICMP is organized.

Chapter Outline

9.1 Introduction

9.2Messages9.3 Debugging Tools

9.4 ICMP Package

9-1 INTRODUCTION

The IP protocol has no error-reporting or error correcting mechanism. What happens if something goes wrong? What happens if a router must discard a datagram because it cannot find a router to the final destination, or because the time-to-live field has a zero value? These are examples of situations where an error has occurred and the IP protocol has no built-in mechanism to notify the original host.

Topics Discussed in the Section

- **✓** The position of ICMP in the TCP/IP suite
- Encapsulation of ICMP Packets

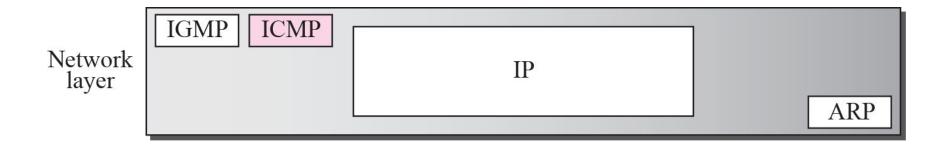
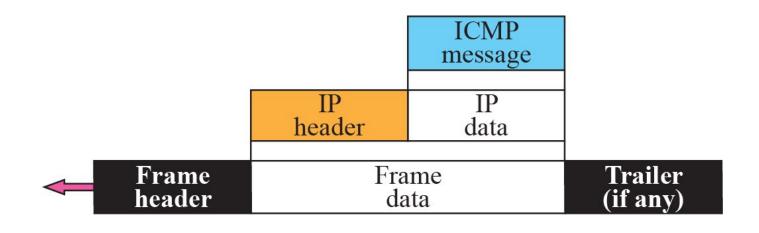


Figure 9.2 ICMP encapsulation



9-2 MESSAGES

ICMP messages are divided into two broad categories: error-reporting messages and query messages. The error-reporting messages report problems that a router or a host (destination) may encounter when it processes an IP packet. The query messages, which occur in pairs, help a host or a network manager get specific information from a router or another host. Also, hosts can discover and learn about routers on their network and routers can help a node redirect its messages.

Topics Discussed in the Section

- **✓** Message Format
- Error Reporting Messages
- Query Messages
- ✓ Checksum

 Table 9.1
 ICMP messages

Category	Туре	Message
	3	Destination unreachable
	4	Source quench
Error-reporting	11	Time exceeded
messages	12	Parameter problem
	5	Redirection
Query	8 or 0	Echo request or reply
messages	13 or 14	Timestamp request or reply

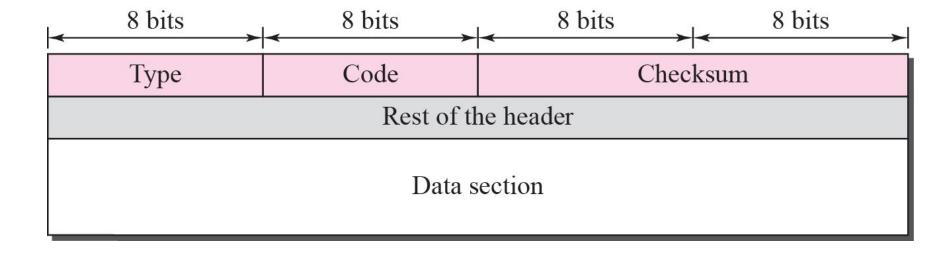
ICMP Parameter Message Format

0 1 2 3 4 5 6 7	8 9 10 11 12 13 14 15	16 17 18 19 20 21 22 23	24 25 26 27 28 29 30 31		
Type	Code	Chec	ksum		
0	1	2	3		
Pointer		Unused			
4	5	6	7		
Internet Header + 8 bytes of Original Data Datagram					
8					

Туре	Code	Meaning	
0	0	Echo Reply	
3	0	Net Unreachable	
8	1	Host Unreachable	
	2	Protocol Unreachable	
	3	Port Unreachable	
	4	Frag needed and DF set	
	5	Source route failed	
	6	Dest network unknown	
	7	Dest host unknown	
	8	Source host isolated	
	9	Network admin prohibited	
	10	Host admin prohibited	
	11	Network unreachable for TOS	
	12	Host unreachable for TOS	
	13	Communication admin prohibited	
4	0	Source Quench (Slow down/Shut up)	

Type	Code	Meaning	
5	0	Redirect datagram for the network	
	1	Redirect datagram for the host	
	2	Redirect datagram for the TOS & Network	
	3	Redirect datagram for the TOS & Host	
8	0	Echo	
9	0	Router advertisement	
10	0	Router selection	
11	0	Time To Live exceeded in transit	
	1	Fragment reassemble time exceeded	
12	0	Pointer indicates the error (Parameter Problem)	
	1	Missing a required option (Parameter Problem)	
	2	Bad length (Parameter Problem)	
13	0	Time Stamp	
14	0	Time Stamp Reply	
15	0	Information Request	
16	0	Information Reply	
17	0	Address Mask Request	
18	0	Address Mask Reply	
30	0	Traceroute (Tracert)	

Figure 9.3 General format of ICMP messages





ICMP always reports error messages to the original source.

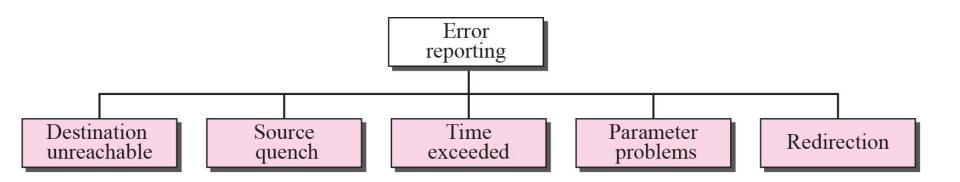


Figure 9.5 Contents of data field for the error message

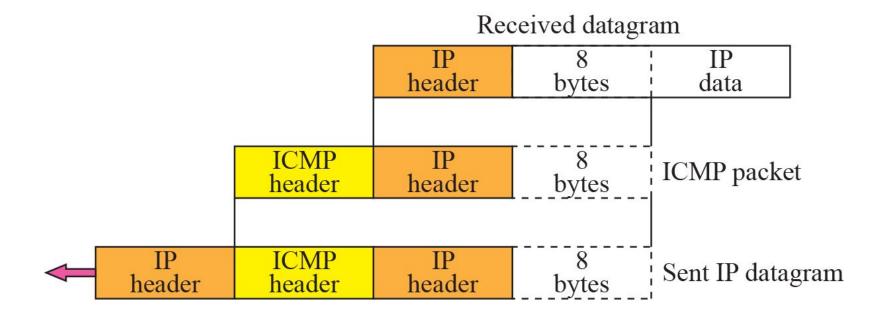
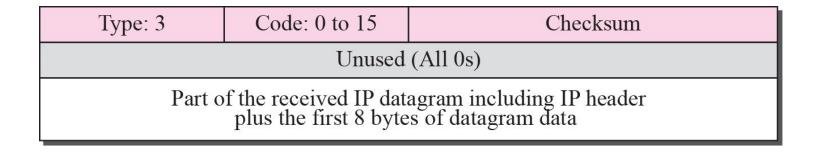


Figure 9.6 Destination-unreachable format



Destination-unreachable messages with codes 2 or 3 can be created only by the destination host.

Other destination-unreachable messages can be created only by routers.

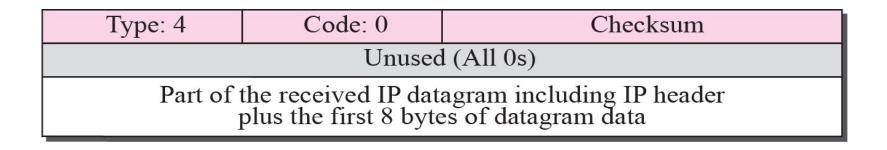


A router cannot detect all problems that prevent the delivery of a packet.



There is no flow-control or congestion-control mechanism in the IP

Figure 9.7 Source-quench format



A source-quench message informs the source that a datagram has been discarded due to congestion in a router or the destination host.

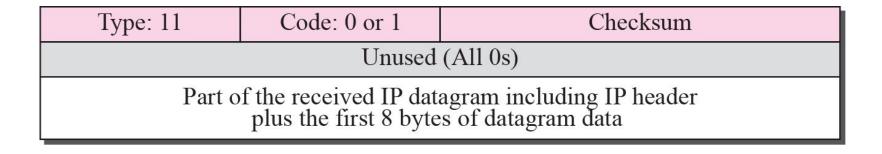
The source must slow down the sending of datagrams until the congestion is relieved.

One source-quench message is sent for each datagram that is discarded due to congestion.

Whenever a router decrements a datagram with a time-to-live value to zero, it discards the datagram and sends a time-exceeded message to the original source.

When the final destination does not receive all of the fragments in a set time, it discards the received fragments and sends a time-exceeded message to the original source.

Figure 9.8 Time-exceeded message format



In a time-exceeded message, code 0 is used only by routers to show that the value of the time-to-live field is zero.

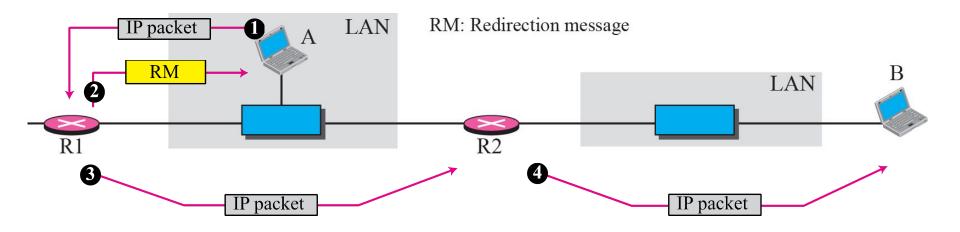
Code 1 is used only by the destination host to show that not all of the fragments have arrived within a set time.

A parameter-problem message can be created by a router or the destination host.

Figure 9.9 Parameter-problem message format

Type: 12	Code: 0 or 1	Checksum
Pointer	Unused (All 0s)	
Part of the received IP datagram including IP header plus the first 8 bytes of datagram data		

Figure 9.10 Redirection concept





A host usually starts with a small routing table that is gradually augmented and updated.

One of the tools to accomplish this is the redirection message.

Figure 9.11 Redirection message format

Type: 5	Code: 0 to 3	Checksum	
IP address of the target router			
Part of the received IP datagram including IP header plus the first 8 bytes of datagram data			

A redirection message is sent from a router to a host on the same local network.

An echo-request message can be sent by a host or router.

An echo-reply message is sent by the host or router that receives an echo-request message.

Echo-request and echo-reply messages can be used by network managers to check the operation of the IP protocol.



Echo-request and echo-reply messages can test the reachability of a host.

This is usually done by invoking the ping command.

Type 8: Echo request Type 0: Echo reply

Type: 8 or 0	Code: 0	Checksum	
Identifier		Sequence number	
Optional data Sent by the request message; repeated by the reply message			

Type 13: request Type 14: reply

Type: 13 or 14	Code: 0	Checksum				
Identifier		Sequence number				
Original timestamp						
Receive timestamp						
Transmit timestamp						

Note

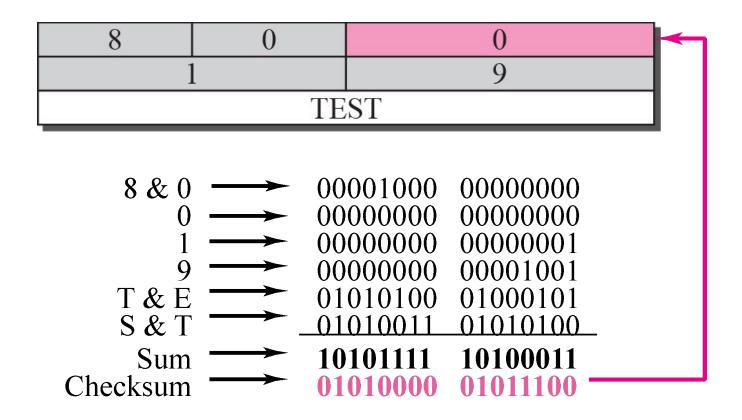
Timestamp-request and timestamp-reply messages can be used to calculate the round-trip time between a source and a destination machine even if their clocks are not synchronized.

Note

The timestamp-request and timestamp-reply messages can be used to synchronize two clocks in two machines if the exact one-way time duration is known.

Figure 9.14 shows an example of checksum calculation for a simple echo-request message (see Figure 9.12). We randomly chose the identifier to be 1 and the sequence number to be 9. The message is divided into 16-bit (2-byte) words. The words are added together and the sum is complemented. Now the sender can put this value in the checksum field.

Figure 9.14 Example of checksum calculation



9-3 DEBUGGING TOOLS

There are several tools that can be used in the Internet for debugging. We can find if a host or router is alive and running. We can trace the route of a packet. We introduce two tools that use ICMP for debugging: ping and traceroute. We will introduce more tools in future chapters after we have discussed the corresponding protocols.

Topics Discussed in the Section

✓ Ping



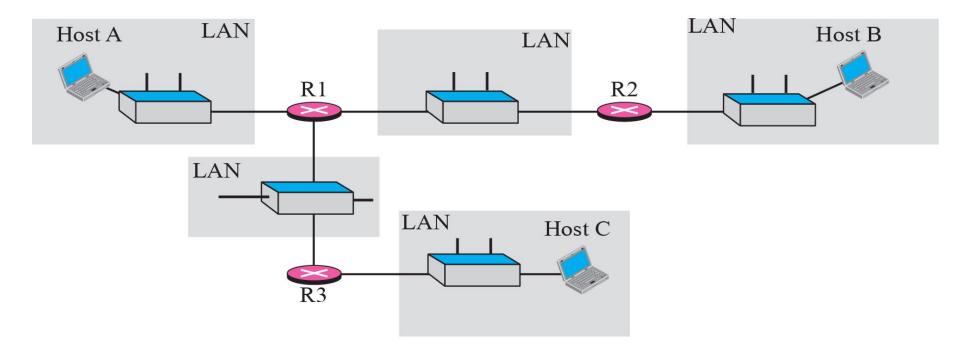
We use the ping program to test the server fhda.edu. The result is shown below:

```
$ ping fhda.edu
PING fhda.edu (153.18.8.1) 56 (84) bytes of data.
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=0
                                                         ttl=62
                                                                   time=1.91 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp seq=1
                                                         ttl=62
                                                                   time=2.04 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=2
                                                         ttl=62
                                                                   time=1.90 ms
                                                                   time=1.97 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp seq=3
                                                         ttl=62
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=4
                                                         ttl=62
                                                                   time=1.93 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=5
                                                         ttl=62
                                                                   time=2.00 \text{ ms}
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=6
                                                         ttl=62
                                                                   time=1.94 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=7
                                                         ttl=62
                                                                   time=1.94 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=8
                                                         ttl=62
                                                                   time=1.97 ms
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=9
                                                                   time=1.89 ms
                                                         ttl=62
64 bytes from tiptoe.fhda.edu (153.18.8.1): icmp_seq=10
                                                                   time=1.98 ms
                                                         ttl=62
--- fhda.edu ping statistics ---
11 packets transmitted, 11 received, 0% packet loss, time 10103 ms
rtt min/avg/max = 1.899/1.955/2.041 ms
```

For the second example, we want to know if the adelphia.net mail server is alive and running. The result is shown below: Note that in this case, we sent 14 packets, but only 13 have been returned. We may have interrupted the program before the last packet, with sequence number 13, was returned.

```
$ ping mail.adelphia.net
PING mail.adelphia.net (68.168.78.100) 56(84) bytes of data.
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=0
                                                             ttl=48
                                                                      time=85.4 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=1
                                                             ttl=48
                                                                      time=84.6 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=2
                                                             ttl=48
                                                                      time=84.9 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=3
                                                             ttl=48
                                                                      time=84.3 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=4
                                                                      time=84.5 ms
                                                             ttl=48
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=5
                                                             ttl=48
                                                                      time=84.7 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=6
                                                             ttl=48
                                                                      time=84.6 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=7
                                                             ttl=48
                                                                      time=84.7 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=8
                                                             ttl=48
                                                                      time=84.4 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=9
                                                             ttl=48
                                                                      time=84.2 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=10
                                                             ttl=48
                                                                      time=84.9 ms
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=11
                                                                      time=84.6 ms
                                                             ttl=48
64 bytes from mail.adelphia.net (68.168.78.100): icmp_seq=12_ttl=48
                                                                      time=84.5 ms
--- mail.adelphia.net ping statistics ---
14 packets transmitted, 13 received, 7% packet loss, time 13129 ms
rtt min/avg/max/mdev = 84.207/84.694/85.469
```

Figure 9.15 The traceroute program operation



We use the traceroute program to find the route from the computer voyager.deanza.edu to the server fhda.edu. The following shows the result.

```
$ traceroute fhda.edu
traceroute to fhda.edu
                            (153.18.8.1), 30 hops max, 38 byte packets
1 Dcore.fhda.edu
                           (153.18.31.25)
                                                0.995 \, \text{ms}
                                                                  0.899 ms
                                                                                   0.878 ms
                                                                  1.064 ms
2 Dbackup.fhda.edu
                           (153.18.251.4)
                                                1.039 ms
                                                                                   1.083 \, \mathrm{ms}
3 tiptoe.fhda.edu
                            (153.18.8.1)
                                                1.797 ms
                                                                  1.642 ms
                                                                                   1.757 ms
```

In this example, we trace a longer route, the route to xerox.com. The following is a partial listing.

<pre>\$ traceroute xerox.com</pre>								
traceroute to xerox.com (13.1.64.93), 30 hops max, 38 byte packets								
1 Dcore.fhda.edu	(153.18.31.254)	0.622 ms	0.891 ms	0.875 ms				
2 Ddmz.fhda.edu	(153.18.251.40)	2.132 ms	2.266 ms	2.094 ms				
3 Cinic.fhda.edu	(153.18.253.126)	2.110 ms	2.145 ms	1.763 ms				
4 cenic.net	(137.164.32.140)	3.069 ms	2.875 ms	2.930 ms				
5 cenic.net	(137.164.22.31)	4.205 ms	4.870 ms	4.197 ms				
6 cenic.net	(137.164.22.167)	4.250 ms	4.159 ms	4.078 ms				
7 cogentco.com	(38.112.6.225)	5.062 ms	4.825 ms	5.020 ms				
8 cogentco.com	(66.28.4.69)	6.070 ms	6.207 ms	5.653 ms				
9 cogentco.com	(66.28.4.94)	6.070 ms	5.928 ms	5.499 ms				

An interesting point is that a host can send a traceroute packet to itself. This can be done by specifying the host as the destination. The packet goes to the loopback address as we expect.

```
$ traceroute voyager.deanza.edu
traceroute to voyager.deanza.edu (127.0.0.1), 30 hops max, 38 byte packets
1 voyager (127.0.0.1) 0.178 ms 0.086 ms 0.055 ms
```

Finally, we use the traceroute program to find the route between fhda.edu and mhhe.com (McGraw-Hill server). We notice that we cannot find the whole route. When traceroute does not receive a response within 5 seconds, it prints an asterisk to signify a problem (not the case in this example), and then tries the next hop.

\$	\$ traceroute mhhe.com							
tra	traceroute to mhhe.com (198.45.24.104), 30 hops max, 38 byte packets							
1	Dcore.fhda.edu	(153.18.31.254)	1.025 ms	0.892 ms	0.880 ms			
2	Ddmz.fhda.edu	(153.18.251.40)	2.141 ms	2.159 ms	2.103 ms			
3	Cinic.fhda.edu	(153.18.253.126)	2.159 ms	2.050 ms	1.992 ms			
4	cenic.net	(137.164.32.140)	3.220 ms	2.929 ms	2.943 ms			
5	cenic.net	(137.164.22.59)	3.217 ms	2.998 ms	2.755 ms			
6	SanJose1.net	(209.247.159.109)	10.653 ms	10.639 ms	10.618 ms			
7	SanJose2.net	(64.159.2.1)	10.804 ms	10.798 ms	10.634 ms			
8	Denver1.Level3.net	(64.159.1.114)	43.404 ms	43.367 ms	43.414 ms			
9	Denver2.Level3.net	(4.68.112.162)	43.533 ms	43.290 ms	43.347 ms			
10	unknown	(64.156.40.134)	55.509 ms	55.462 ms	55.647 ms			
11	mcleodusa1.net	(64.198.100.2)	60.961 ms	55.681 ms	55.461 ms			
12	mcleodusa2.net	(64.198.101.202)	55.692 ms	55.617 ms	55.505 ms			
13	mcleodusa3.net	(64.198.101.142)	56.059 ms	55.623 ms	56.333 ms			
14	mcleodusa4.net	(209.253.101.178)	297.199 ms	192.790 ms	250.594 ms			
15	eppg.com	(198.45.24.246)	71.213 ms	70.536 ms	70.663 ms			
16					•••			

9-4 ICMP PACKAGE

To give an idea of how ICMP can handle the sending and receiving of ICMP messages, we present our version of an ICMP package made of two modules: an input module and an output module.

Topics Discussed in the Section

- **✓** Input Module
- **Output Module**



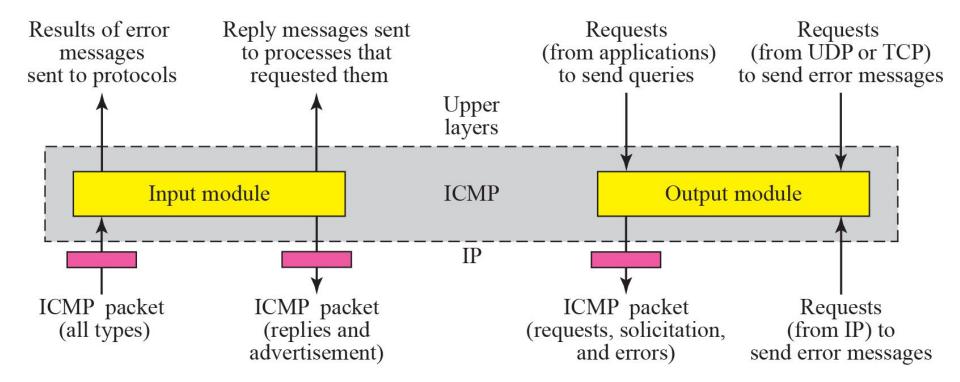


 Table 9.2
 Input Module

```
ICMP_Input_module (ICMP_Packet)
 2
    {
 3
       If (the type is a request)
 4
       {
            Create a reply
            Send the reply
        }
       If (the type defines a redirection)
 8
        {
 9
            Modify the routing table
10
11
       }
12
       If (the type defines other error messages)
13
        {
14
            Inform the appropriate source protocol
       }
15
16
       Return
17
```

 Table 9.3
 Output Module

```
1
    ICMP_Output_Module (demand)
 2
    {
 3
       If (the demand defines an error message)
 4
 5
                If (demand comes from IP AND is forbidden)
 6
                {
 7
                     Return
 8
 9
                If (demand is a valid redirection message)
10
11
                     Return
12
13
                Create an error message
14
       If (demand defines a request)
15
        {
16
             Create a request message
17
18
       Send the message
19
       Return
20
    }
```